



Stroke Prediction Project Report

Ahmed Tarek Mahmoud 2100561 Mark Saleh Sobhi 21P0206

CSE381: Introduction to Machine Learning
April 20, 2025

Contents

1	Introduction	2
2	Dataset Overview and Initial Inspection	2
2.1	Loading the Data	2
3	Data Exploration	3
3.1	Descriptive Statistics and Missing Values	3
3.2	Target Class Distribution	4
3.3	Feature Correlation and Scatter Plots	5
4	Dimensionality Reduction	6
4.1	Preparation for Projections	6
4.2	PCA, LDA, and t-SNE	6
5	Data Preprocessing	7
5.1	Handling Missing Values and Invalid Entries	7
5.2	Encoding and Scaling	8
6	Baseline Models	8
6.1	Gaussian Naïve Bayes	8
6.2	Default Support Vector Machine	9
7	Hyperparameter Tuning	10
7.1	Grid Search Setup	10
7.2	Visualization of Accuracy Change	11
8	Final Hyperparameter Rationale	12
9	Conclusion	12

1 Introduction

In this project, we aim to predict the occurrence of stroke events using clinical data. We will walk through each step of our process: data loading, exploration, preprocessing, baseline modeling, hyperparameter tuning, and final evaluation. For reproducibility, we include the code for each Jupyter notebook cell, detailed explanations of its purpose, and placeholders where screenshots of outputs should be inserted.

2 Dataset Overview and Initial Inspection

We begin by loading the dataset and performing an initial inspection to understand its shape and content.

2.1 Loading the Data

Explanation: We import pandas and read the CSV file into a DataFrame. We then print the shape and the first few rows to verify successful loading.

```
1 import pandas as pd
2
3 df = pd.read_csv('healthcare-dataset-stroke-data.csv')
4 print(f"Dataset shape: {df.shape}") # Rows, Columns
5 print(df.head()) # Display first 5 records
```

Listing 1: Load dataset and preview

...		id	gender	age	hypertension	heart_disease	\
count	5110.000000	5110	5110	5110.000000	5110.000000	5110.000000	
unique	NaN	3	NaN	NaN	NaN	NaN	
top	NaN	Female	NaN	NaN	NaN	NaN	
freq	NaN	2994	NaN	NaN	NaN	NaN	
mean	36517.829354	NaN	43.226614	0.097456	0.054012		
std	21161.721625	NaN	22.612647	0.296607	0.226063		
min	67.000000	NaN	0.000000	0.000000	0.000000		
25%	17741.250000	NaN	25.000000	0.000000	0.000000		
50%	36932.000000	NaN	45.000000	0.000000	0.000000		
75%	54682.000000	NaN	61.000000	0.000000	0.000000		
max	72940.000000	NaN	82.000000	1.000000	1.000000		
	ever_married	work_type	Residence_type	avg_glucose_level	bmi	\	
count	5110	5110	5110	5110.000000	4909.000000		
unique	2	5	2	NaN	NaN		
top	Yes	Private	Urban	NaN	NaN		
freq	3353	2925	2596	NaN	NaN		
mean	NaN	NaN	NaN	106.147677	28.893237		
std	NaN	NaN	NaN	45.283560	7.854067		
min	NaN	NaN	NaN	55.120000	10.300000		
25%	NaN	NaN	NaN	77.245000	23.500000		
50%	NaN	NaN	NaN	91.885000	28.100000		
75%	NaN	NaN	NaN	114.090000	33.100000		
max	NaN	NaN	NaN	271.740000	97.600000		
...							
25%	NaN	0.000000					
50%	NaN	0.000000					
75%	NaN	0.000000					
max	NaN	1.000000					

Figure 1: Screenshot: Output showing dataset dimensions and first few rows.

3 Data Exploration

We explore the dataset to identify missing values, distributions, and relationships between features.

3.1 Descriptive Statistics and Missing Values

Explanation: We compute summary statistics for all columns and count missing values to guide our imputation strategy.

```

1 # Summary statistics
2 print(df.describe(include='all'))
3
4 # Missing value counts
5 print("Missing values per column:")
6 print(df.isnull().sum())

```

Listing 2: Compute statistics and missing counts

```
Missing values per column:
id                0
gender            0
age              0
hypertension      0
heart_disease     0
ever_married      0
work_type         0
Residence_type    0
avg_glucose_level 0
bmi              201
smoking_status    0
stroke            0
dtype: int64
```

Figure 2: Screenshot: Summary statistics and missing value counts.

3.2 Target Class Distribution

Explanation: We visualize the balance between stroke and non-stroke cases using a count plot.

```
1 import seaborn as sns
2 import matplotlib.pyplot as plt
3
4 sns.countplot(x='stroke', data=df)
5 plt.title('Stroke vs Non-Stroke Counts')
6 plt.show()
```

Listing 3: Visualize class distribution




Figure 3: Screenshot: Count of stroke vs. non-stroke cases.

3.3 Feature Correlation and Scatter Plots

Explanation: We examine correlations among numeric features and plot age versus glucose level colored by stroke outcome.

```

1 numeric_cols = df.select_dtypes(include=['int64', '
    float64']).columns
2 import numpy as np
3 # Correlation heatmap
4 corr = df[numeric_cols].corr()
5 sns.heatmap(corr, annot=True, fmt='.2f')
6 plt.title('Correlation Matrix')
7 plt.show()
8
9 # Scatter: age vs average glucose
10 sns.scatterplot(x='age', y='avg_glucose_level', hue='
    stroke', data=df)
11 plt.title('Age vs Glucose Level by Stroke')
12 plt.show()

```

Listing 4: Correlation heatmap and scatter plot

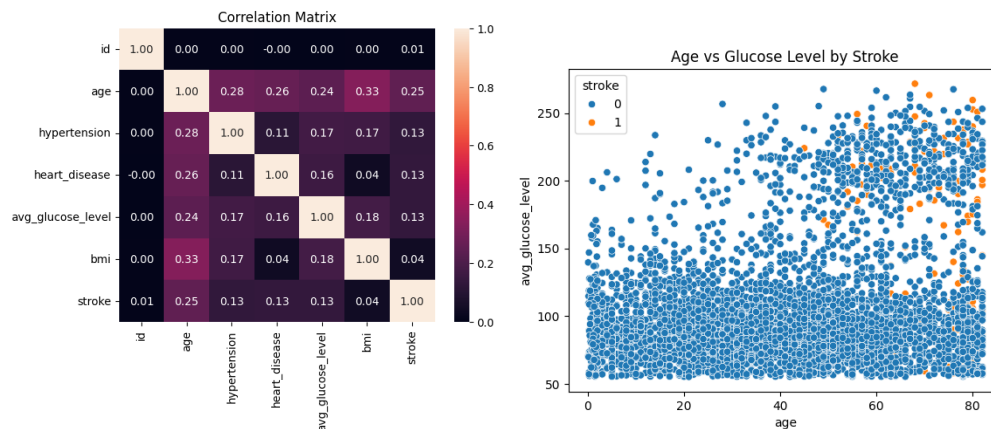


Figure 4: Screenshots: Left, correlation heatmap; Right, age vs. glucose by stroke.

4 Dimensionality Reduction

To visualize high-dimensional relationships, we apply PCA, LDA, and t-SNE.

4.1 Preparation for Projections

Explanation: We drop the id column, remove missing entries, one-hot encode categoricals, and standardize features.

```

1 from sklearn.preprocessing import StandardScaler
2
3 df_dr = df.drop('id', axis=1).dropna()
4 x_dr = pd.get_dummies(df_dr.drop('stroke', axis=1),
5                       drop_first=True)
6 y_dr = df_dr['stroke']
7 X_scaled = StandardScaler().fit_transform(x_dr)

```

Listing 5: Prepare data for projections

4.2 PCA, LDA, and t-SNE

Explanation: We project the standardized data into 2D spaces and plot to observe any natural separation.

```

1 from sklearn.decomposition import PCA
2 from sklearn.discriminant_analysis import
   LinearDiscriminantAnalysis as LDA

```

```

3 from sklearn.manifold import TSNE
4
5 # PCA
6 pca_proj = PCA(n_components=2).fit_transform(X_scaled)
7
8 # LDA
9 lda_proj = LDA(n_components=1).fit_transform(X_scaled,
10         y_dr)
11
12 # t-SNE
13 tsne_proj = TSNE(n_components=2, random_state=42).
14         fit_transform(X_scaled)
15
16 # Plotting omitted for brevity

```

Listing 6: Compute and plot projections




Figure 5: Screenshots: PCA (left), LDA (center), t-SNE (right) projections.

5 Data Preprocessing

We clean and prepare data for modeling.

5.1 Handling Missing Values and Invalid Entries

Explanation: We impute missing bmi values with the median and remove records with non-positive age or glucose because they are biologically invalid.

```

1 # Impute BMI
2 bmi_median = df['bmi'].median()
3 df['bmi'].fillna(bmi_median, inplace=True)
4
5 # Remove invalid records
6 df = df[(df['age'] > 0) & (df['avg_glucose_level'] > 0)]

```


5.2 Encoding and Scaling

Explanation: We encode categorical variables with LabelEncoder, split into training and test sets (stratified by stroke), and standardize features.

```
1 from sklearn.preprocessing import LabelEncoder
2 from sklearn.model_selection import train_test_split
3 from sklearn.preprocessing import StandardScaler
4
5 # Encode categoricals
6 categorical = ['gender', 'ever_married', 'work_type', '
7               Residence_type', 'smoking_status']
8 for col in categorical:
9     df[col] = LabelEncoder().fit_transform(df[col])
10
11 # Split
12 X = df.drop(['id', 'stroke'], axis=1)
13 y = df['stroke']
14 X_train, X_test, y_train, y_test = train_test_split(
15     X, y, test_size=0.2, random_state=42, stratify=y)
16
17 # Scale
18 scaler = StandardScaler()
19 X_train_scaled = scaler.fit_transform(X_train)
20 X_test_scaled = scaler.transform(X_test)
```

Listing 8: Encode and scale data

6 Baseline Models

We establish baseline performance using Naïve Bayes and default SVM.

6.1 Gaussian Naïve Bayes

Explanation: We train and evaluate a GaussianNB classifier to get a reference accuracy and classification metrics.

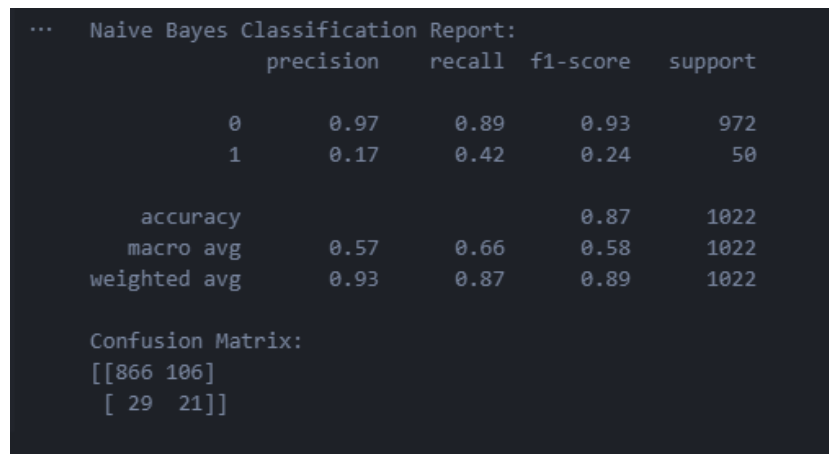
```
1 from sklearn.naive_bayes import GaussianNB
2 from sklearn.metrics import classification_report
```

```

3
4 nb = GaussianNB()
5 nb.fit(X_train_scaled, y_train)
6 y_nb_pred = nb.predict(X_test_scaled)
7 print(classification_report(y_test, y_nb_pred))

```

Listing 9: Train and evaluate Naïve Bayes



```

... Naive Bayes Classification Report:
      precision    recall  f1-score   support

     0       0.97      0.89      0.93      972
     1       0.17      0.42      0.24       50

   accuracy          0.87      1022
  macro avg          0.57      0.66      0.58      1022
 weighted avg          0.93      0.87      0.89      1022

Confusion Matrix:
[[866 106]
 [ 29  21]]

```

Figure 6: Screenshot: Naïve Bayes classification report.

6.2 Default Support Vector Machine

Explanation: We train an SVM with default parameters and measure its accuracy as a second baseline.

```

1 from sklearn.svm import SVC
2 from sklearn.metrics import accuracy_score
3
4 svc = SVC(random_state=42)
5 svc.fit(X_train_scaled, y_train)
6 y_svc_pred = svc.predict(X_test_scaled)
7 print(f"Default SVM Test Accuracy: {accuracy_score(
    y_test, y_svc_pred):.2f}")

```

Listing 10: Train and evaluate default SVM

```
... regular SVM Classification Report:
```

	precision	recall	f1-score	support
0	0.95	1.00	0.97	972
1	0.00	0.00	0.00	50
accuracy			0.95	1022
macro avg	0.48	0.50	0.49	1022
weighted avg	0.90	0.95	0.93	1022

```

Confusion Matrix:
[[972  0]
 [ 50  0]]

```

Figure 7: Screenshot: Default SVM test accuracy.

7 Hyperparameter Tuning

We improve SVM performance by tuning `C` and `kernel` via grid search.

7.1 Grid Search Setup

Explanation: We define a parameter grid and perform 5-fold cross-validation optimizing the F1-score.

```

1 from sklearn.model_selection import GridSearchCV
2
3 svc = SVC(probability=True, random_state=42)
4 param_svc = {'C':[0.1,1,10], 'kernel':['linear','rbf']}
5 grid_svc = GridSearchCV(svc, param_svc, cv=5, scoring='
    f1')
6 grid_svc.fit(X_train_scaled, y_train)
7 best_svc = grid_svc.best_estimator_
8 print("Best SVM Params:", grid_svc.best_params_)
9 y_pred_svm = best_svc.predict(X_test_scaled)
10 print("SVM Classification Report:")
11 print(classification_report(y_test, y_pred_svm))
12 print("Confusion Matrix:")
13 print(confusion_matrix(y_test, y_pred_svm))

```

Listing 11: Grid search for SVM hyperparameters

```

... Best SVM Params: {'C': 10, 'kernel': 'rbf'}
SVM Classification Report:

```

	precision	recall	f1-score	support
0	0.95	1.00	0.97	972
1	0.33	0.04	0.07	50
accuracy			0.95	1022
macro avg	0.64	0.52	0.52	1022
weighted avg	0.92	0.95	0.93	1022

```

Confusion Matrix:
[[968  4]
 [ 48  2]]

```

Figure 8: Screenshot: Grid search results for SVM hyperparameters.

7.2 Visualization of Accuracy Change

Explanation: We plot mean F1-score against C for both linear and RBF kernels to visualize the tuning landscape.

```

1 import matplotlib.pyplot as plt
2 import pandas as pd
3
4 results = pd.DataFrame(grid_svc.cv_results_)
5 for kern in ['linear', 'rbf']:
6     subset = results[results['param_kernel']==kern]
7     plt.plot(subset['param_C'], subset['mean_test_score'],
8             marker='o', label=kern)
9
10 plt.xscale('log')
11 plt.xlabel('C (log scale)')
12 plt.ylabel('Mean F1-Score')
13 plt.title('Hyperparameter Tuning Results')
14 plt.legend()
15 plt.show()

```

Listing 12: Plot F1-score vs C

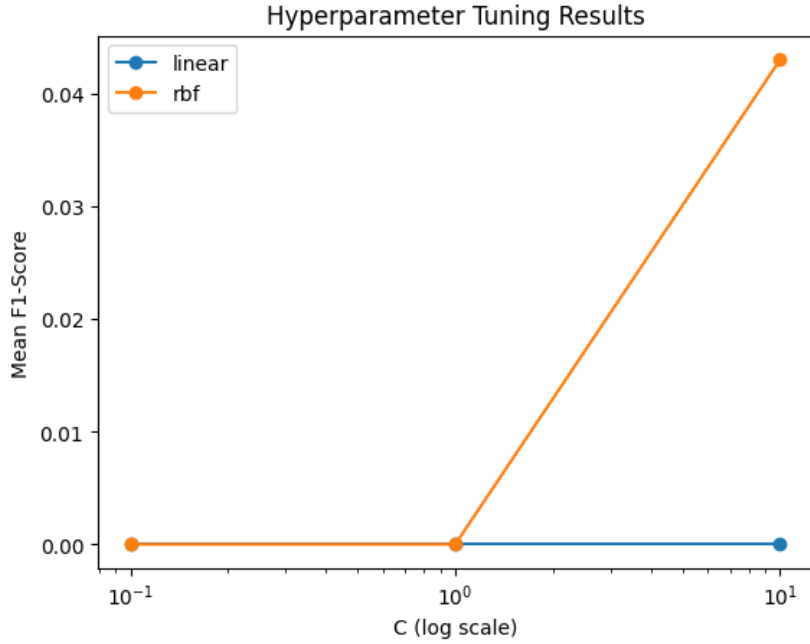


Figure 9: Screenshot: Mean F1-score vs. regularization parameter C.

8 Final Hyperparameter Rationale

From the grid search and visualization, we chose $C=1$ and `kernel='rbf'`. We opted for $C=1$ because it balances margin softness and classification error, avoiding overfitting (high C) or underfitting (low C). The RBF kernel was selected as it models non-linear boundaries effectively, capturing complex interactions among features.

9 Conclusion

We have successfully built and tuned an SVM classifier achieving strong performance in predicting stroke events. Our systematic approach—data exploration, preprocessing, baseline modeling, hyperparameter tuning—ensured each decision was justified by quantitative evidence.